

DATA traffic analysis

This notebook contains analysis about the data traffic generated by three different kinds of source:

- * Download a page from as.com (www.as.com)
- * Download a video from youtube (https://www.youtube.com/watch?v=ruabyha_mGI)
- * Simulate voIP packets (
 - * Bits per second
CODEC=64000
 - * Voice payload seconds per packet
PACKETVOICE=0.02
 - * 8 bits are one byte
BIT_TO_BYTE=8
 - IPHEADER=40 # BYTES
 - ETHERNETHEADER=18 # BYTES
 - VOIP_PACKET_SIZE=\$(echo "(\$CODEC * \$PACKETVOICE + \$IPHEADER + \$ETHERNETHEADER) * (1/\$BIT_TO_BYTE)" | bc -l | awk '{print int(\$1+0.5)}'))

```
In [1]: import pandas as pd
import os
import shutil
from matplotlib import pyplot as plt
import seaborn as sns
import numpy as np
%matplotlib inline
```

```
In [2]: data_name = 'data_traffic'
path_to_data = '../Data/'
sep = '#'
file = path_to_data + data_name
TFM_all_data = pd.read_csv(file, sep=sep)
```

```
In [3]: TFM_all_data.iloc[0:4]
```

```
Out[3]:
```

| | is_youtube | IP_FiveTuple | IP_SrcIP | proto | timeStamp | coord_1 | dpiPktNum | IP_UpLink | IP_TotLen | IP_DstIP |
|---|------------|---|---------------|-------|-----------------|---------|-----------|-----------|-----------|---------------|
| 0 | 0 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 192.168.1.135 | tcp | 20:13:13.499324 | 10 | 1-7 | 1 | 0 | 2.16.8.43 |
| 1 | 0 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 2.16.8.43 | tcp | 20:13:13.513781 | 71 | 0-8 | 0 | 0 | 192.168.1.135 |
| 2 | 0 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 192.168.1.135 | tcp | 20:13:13.513812 | 72 | 1-9 | 1 | 0 | 2.16.8.43 |
| 3 | 0 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 192.168.1.135 | tcp | 20:13:13.513914 | 22 | 1-10 | 1 | 197 | 2.16.8.43 |

We will mix traffic labels into one column. First, we will set label name instead of 1 and 0 of each label. Then we remove each label column.

```
In [4]: # Mix labels
video_label = 'is_youtube'
voip_label = 'voIP'
brow_label = 'browsing'
TFM_data_df = TFM_all_data.copy()
TFM_data_df.loc[TFM_data_df[video_label] == 1, 'label'] = video_label
TFM_data_df.loc[TFM_data_df[voip_label] == 1, 'label'] = voip_label
TFM_data_df.loc[TFM_data_df[brow_label] == 1, 'label'] = brow_label

# Remove each label column
column_names = TFM_data_df.columns.tolist()
column_names.remove(video_label)
column_names.remove(voip_label)
column_names.remove(brow_label)
TFM_data_df = TFM_data_df[column_names]

TFM_data_df.iloc[0:4]
```

Out[4]:

| | IP_FiveTuple | IP_SrcIP | proto | timeStamp | coord_1 | dpiPktNum | IP_UpLink | IP_TotLen | IP_DstIP | IP_Versior |
|---|---|---------------|-------|-----------------|---------|-----------|-----------|-----------|---------------|------------|
| 0 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 192.168.1.135 | tcp | 20:13:13.499324 | 10 | 1-7 | 1 | 0 | 2.16.8.43 | 4 |
| 1 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 2.16.8.43 | tcp | 20:13:13.513781 | 71 | 0-8 | 0 | 0 | 192.168.1.135 | 4 |
| 2 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 192.168.1.135 | tcp | 20:13:13.513812 | 72 | 1-9 | 1 | 0 | 2.16.8.43 | 4 |
| 3 | 17;192.168.1.135/0;2.16.8.43/0;80;46710 | 192.168.1.135 | tcp | 20:13:13.513914 | 22 | 1-10 | 1 | 197 | 2.16.8.43 | 4 |

Flows

Following some read papers, main data of a flow is contained in the first 4.5 seconds of a flow of packets. According with this data, we are going to calculate (helped by timeStamp column) the number of packets in 4.5 seconds for each label.

```
In [5]: vid_df = TFM_data_df.loc[TFM_data_df['label'] == video_label].reset_index(drop=True)
bro_df = TFM_data_df.loc[TFM_data_df['label'] == brow_label].reset_index(drop=True)
voip_df = TFM_data_df.loc[TFM_data_df['label'] == voip_label].reset_index(drop=True)
```

```

In [6]: init_vid = vid_df['timeStamp'][0]
init_bro = bro_df['timeStamp'][0]
init_voip = voip_df['timeStamp'][0]

{'video_init_hour': init_vid, 'voip_init_hour': init_voip, 'bro_init_hour': init_bro}

Out[6]: {'bro_init_hour': '20:13:13.499324',
'video_init_hour': '20:34:12.451696',
'voip_init_hour': '21:22:07.297557'}

In [7]: vid_num_packets = 2100
print('HORA INICIO VIDEO: ' + str(init_vid) + 'HORA CALCULADA: ' + str(vid_df.iloc[vid_num_packets][3]))
HORA INICIO VIDEO: 20:34:12.451696HORA CALCULADA: 20:34:16.305427

In [8]: bro_num_packets = 350
print('HORA INICIO BROWSING: ' + str(init_bro) + 'HORA CALCULADA: ' + str(bro_df.iloc[bro_num_packets][3]))
HORA INICIO BROWSING: 20:13:13.499324HORA CALCULADA: 20:13:17.887830

In [9]: voip_num_packets = 400
print('HORA INICIO VOIP: ' + str(init_voip) + 'HORA CALCULADA: ' + str(voip_df.iloc[voip_num_packets][3]))
HORA INICIO VOIP: 21:22:07.297557HORA CALCULADA: 21:22:11.385689

In [10]: (voip_num_packets + bro_num_packets + vid_num_packets) / 3
Out[10]: 950.0

```

Important INFO

Now, for the creation of the model, we will remove some columns that depend on where the demo is being running. The goal is create a demo that can be showed in any place. Due to that, we will remove:

- * TimeStamp: Only needed for the visualization in Kibana.
- * coord_1 and coord_2: right now are created in random way so are not important.
- * IP_Version: always version 4
- * dpiPktNum: number of packet generated. Is a number set by the generator so not important here.

IP_SrcIP, IP_DstIP and IP_FiveTuple contains info about IPs and ports of the communication. It could be relevant information but that tie us to have a concrete IP in the device where the data traffic is being analyzed.

```
In [11]: columns_to_remove = ['IP_FiveTuple', 'IP_SrcIP', 'timeStamp', 'coord_1',  
                             'dpiPktNum', 'IP_UpLink', 'IP_DstIP', 'IP_Version', 'coord_2']  
for col in columns_to_remove:  
    column_names.remove(col)  
  
TFM_data_df = TFM_data_df[column_names]  
TFM_data_df = TFM_data_df[['label', 'proto', 'IP_TotLen']]  
df_rows = TFM_data_df.shape[0]  
TFM_data_df.iloc[0:4]
```

Out[11]:

| | label | proto | IP_TotLen |
|---|----------|-------|-----------|
| 0 | browsing | tcp | 0 |
| 1 | browsing | tcp | 0 |
| 2 | browsing | tcp | 0 |
| 3 | browsing | tcp | 197 |

```
In [12]: TFM_data_df.groupby('label').count()
```

Out[12]:

| | proto | IP_TotLen |
|------------|-------|-----------|
| label | | |
| browsing | 38112 | 38112 |
| is_youtube | 56235 | 56235 |
| voIP | 52183 | 52183 |

```
In [13]: TFM_data_df.groupby('label').describe()
```

Out[13]:

| | IP_TotLen | | | | | | | |
|------------|-----------|-------------|-------------|-----|-------|--------|--------|---------|
| | count | mean | std | min | 25% | 50% | 75% | max |
| label | | | | | | | | |
| browsing | 38112.0 | 2429.369674 | 3497.168700 | 0.0 | 0.0 | 133.0 | 4344.0 | 27512.0 |
| is_youtube | 56235.0 | 1757.337619 | 1989.861707 | 0.0 | 0.0 | 1448.0 | 2896.0 | 17376.0 |
| voIP | 52183.0 | 175.132629 | 27.367425 | 0.0 | 175.0 | 175.0 | 175.0 | 4155.0 |

```
In [14]: TFM_data_df.groupby('proto').count()
```

Out[14]:

| | label | IP_TotLen |
|-------|-------|-----------|
| proto | | |
| ICMP | 51956 | 51956 |
| UDP | 1508 | 1508 |
| tcp | 93066 | 93066 |

```
In [15]: TFM_data_df.groupby('proto').describe()
```

Out[15]:

| | IP_TotLen | | | | | | | |
|-------|-----------|-------------|-------------|-------|-------|--------|--------|---------|
| | count | mean | std | min | 25% | 50% | 75% | max |
| proto | | | | | | | | |
| ICMP | 51956.0 | 175.000000 | 0.000000 | 175.0 | 175.0 | 175.0 | 175.0 | 175.0 |
| UDP | 1508.0 | 83.465517 | 85.238964 | 24.0 | 28.0 | 35.0 | 133.0 | 451.0 |
| tcp | 93066.0 | 2055.882900 | 2730.843771 | 0.0 | 0.0 | 1448.0 | 2896.0 | 27512.0 |

```
In [16]: TFM_data_df.groupby(['label', 'proto']).count()
```

Out[16]:

| IP_TotLen | | |
|------------|-------|-------|
| label | proto | |
| browsing | UDP | 1293 |
| | tcp | 36819 |
| is_youtube | UDP | 69 |
| | tcp | 56166 |
| voIP | ICMP | 51956 |
| | UDP | 146 |
| | tcp | 81 |

```
In [17]: TFM_data_df.groupby(['label', 'proto']).describe()
```

Out[17]:

| IP_TotLen | | | | | | | | | |
|------------|-------|---------|-------------|-------------|-------|-------|--------|--------|---------|
| | | count | mean | std | min | 25% | 50% | 75% | max |
| label | proto | | | | | | | | |
| browsing | UDP | 1293.0 | 82.337974 | 87.083690 | 24.0 | 28.0 | 56.0 | 133.0 | 451.0 |
| | tcp | 36819.0 | 2511.792118 | 3529.756939 | 0.0 | 0.0 | 226.0 | 4344.0 | 27512.0 |
| is_youtube | UDP | 69.0 | 124.275362 | 84.941136 | 33.0 | 42.0 | 83.0 | 213.0 | 246.0 |
| | tcp | 56166.0 | 1759.343838 | 1990.257499 | 0.0 | 0.0 | 1448.0 | 2896.0 | 17376.0 |
| voIP | ICMP | 51956.0 | 175.000000 | 0.000000 | 175.0 | 175.0 | 175.0 | 175.0 | 175.0 |
| | UDP | 146.0 | 74.164384 | 60.549214 | 31.0 | 35.0 | 35.0 | 154.0 | 300.0 |
| | tcp | 81.0 | 442.197531 | 625.349131 | 0.0 | 0.0 | 257.0 | 1033.0 | 4155.0 |

Dataframe transformations

In order to plot and to prepare data for fitting, we will apply some transformations

```
In [18]: TFM_data_model = TFM_data_df.copy()
proto_keys = {'tcp': 1, 'UDP': 2, 'ICMP': 3}
label_keys = {video_label: 1, voip_label: 2, brow_label: 3}

TFM_data_model.loc[TFM_data_model['proto'] == 'tcp', 'proto'] = proto_keys['tcp']
TFM_data_model.loc[TFM_data_model['proto'] == 'UDP', 'proto'] = proto_keys['UDP']
TFM_data_model.loc[TFM_data_model['proto'] == 'ICMP', 'proto'] = proto_keys['ICMP']

TFM_data_model.loc[TFM_data_model['label'] == video_label, 'label'] = label_keys[video_label]
TFM_data_model.loc[TFM_data_model['label'] == voip_label, 'label'] = label_keys[voip_label]
TFM_data_model.loc[TFM_data_model['label'] == brow_label, 'label'] = label_keys[brow_label]
```

```
In [19]: def to_csv(path_tocsv, file_name_tocsv, df, sep = '#', header = True):
        data_path_tocsv = path_tocsv + file_name_tocsv
        df.to_csv(path_or_buf= data_path_tocsv, sep=sep, header = True, index=False)
```

```
In [20]: path_tocsv = '../Data/'
file_name_tocsv = 'data_model'
df = TFM_data_model
to_csv(path_tocsv, file_name_tocsv, df)
```

PLOTS


According to some read papers, only with packet length, the source of the data traffic could be guessed. Due to that, we will make some plots. We plot the flow of datatrafic for each label and for each protocol

```
In [21]: vid_df = TFM_data_df.loc[TFM_data_df['label'] == video_label, 'IP_TotLen'].tolist()[0:vid_num_packets]
bro_df = TFM_data_df.loc[TFM_data_df['label'] == brow_label, 'IP_TotLen'].tolist()[0:bro_num_packets]
voip_df = TFM_data_df.loc[TFM_data_df['label'] == voip_label, 'IP_TotLen'].tolist()[0:voip_num_packets]
```




```
In [22]: plt.plot(voip_df)
```

```
Out[22]: [<matplotlib.lines.Line2D at 0x7fe66e6ba898>]
```




```
In [23]: plt.plot(bro_df)
```

```
Out[23]: [<matplotlib.lines.Line2D at 0x7fe671562278>]
```




```
In [24]: plt.plot(vid_df)
```

```
Out[24]: [<matplotlib.lines.Line2D at 0x7fe670dafc88>]
```




```
In [25]: # Plot grouping by protocol
plt.scatter(TFM_data_model['proto'].tolist(),
            TFM_data_model['IP_TotLen'].tolist(),
            color=['red', 'green', 'blue'])

plt.show()
```



```
In [26]: # Plot grouping by label
plt.scatter(TFM_data_model['label'].tolist(),
            TFM_data_model['IP_TotLen'].tolist(),
            color=['red', 'green', 'blue'])

plt.show()
```



Fitting

```
In [1]: # Spark
import findspark
spark_path = "/home/nacho/spark"
findspark.init(spark_path)
import pyspark
from pyspark.sql import SparkSession
spark = (SparkSession.builder
        .master("local[*]")
        .config("spark.driver.cores", 1)
        .appName("understanding_sparksession")
        .getOrCreate() )
sc = spark.sparkContext

print(spark)
print(sc)
from pyspark.sql.types import *
from pyspark.sql.functions import *
from pyspark.sql.functions import col, when
from pyspark.ml import Pipeline
from pyspark.ml.feature import StringIndexer
from pyspark.ml import Pipeline
# RANDOM FOREST
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.feature import IndexToString, StringIndexer, VectorIndexer
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.classification import RandomForestClassificationModel
# GRADIENT BOOSTED TREE
from pyspark.ml.classification import GBClassifier
from pyspark.ml.feature import StringIndexer, VectorIndexer
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
# MULTILAYER PERCEPTRON
from pyspark.ml.classification import MultilayerPerceptronClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

<pyspark.sql.session.SparkSession object at 0x7f177765da20>
<pyspark.context.SparkContext object at 0x7f17a7720400>
```

From CSV to LIBSVM format (only for numerical data)

For all models in spark is needed that training files are in *libsvm* format. Due to that, we will need to use the next functions


```
In [28]: import sys
import csv
from collections import defaultdict

# Source: https://github.com/zygmuntz/phraug/blob/master/csv2libsvm.py

def construct_line( label, line ):
    new_line = []
    if float( label ) == 0.0:
        label = "0"
    new_line.append( label )

    for i, item in enumerate( line ):
        if item == ':' or float( item ) == 0.0:
            continue
        new_item = "%s:%s" % ( i + 1, item )
        new_line.append( new_item )
    new_line = " ".join( new_line )
    new_line += "\n"
    return new_line

def csv2libsvm(inputfile, outoutfile, labels=0,headers=0):
    """
    Convert CSV file to libsvm format. Works only with numeric variables.
    Put -1 as label index (label) if there are no labels in your file.
    Expecting no headers. If present, headers can be skipped with headers == 1.

    Convert CSV to the LIBSVM format. If there are no labels in the input file,
    specify label index = -1. If there are headers in the input file, specify skip headers = 1.
    """

    input_file = inputfile
    output_file = outoutfile

    try:
        label_index = int( labels )
    except IndexError:
        label_index = 0

    try:
        skip_headers = headers
```

```

except IndexError:
    skip_headers = 0

i = open( input_file, 'r' )
o = open( output_file, 'w' )

reader = csv.reader( i , delimiter = '#' )

if skip_headers:
    headers = next(reader)

for line in reader:
    if label_index == -1:
        label = '1'
    else:
        label = line.pop( label_index )

    new_line = construct_line( label, line )

    o.write( new_line )

```

We transform the desired data into the correct format

```

In [29]: read_path = '../Data/'
input_file = 'data_model'
output_file = 'data_model_libsvm.csv'

csv_input = read_path + input_file
csv_output = read_path + output_file
csv2libsvm(csv_input, csv_output, labels=0, headers=1)

```

Algorithms

Gradient-boosted tree classifier

```
In [32]: from pyspark.ml import Pipeline
from pyspark.ml.classification import GBTClassifier
from pyspark.ml.feature import StringIndexer, VectorIndexer
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

GBTmodel_path = "../Model/Gradient_boosted_tree"

def gradient_boosted_tree(data, model_path = GBTmodel_path):

    # Index labels, adding metadata to the label column.
    # Fit on whole dataset to include all labels in index.
    labelIndexer = StringIndexer(inputCol="label", outputCol="indexedLabel").fit(data)
    # Automatically identify categorical features, and index them.
    # Set maxCategories so features with > 4 distinct values are treated as continuous.
    featureIndexer = \
        VectorIndexer(inputCol="features", outputCol="indexedFeatures", maxCategories=4).fit(data)

    # Split the data into training and test sets (30% held out for testing)
    (trainingData, testData) = data.randomSplit([0.7, 0.3])

    # Train a GBT model.
    gbt = GBTClassifier(labelCol="indexedLabel", featuresCol="indexedFeatures", maxIter=10)

    # Chain indexers and GBT in a Pipeline
    pipeline = Pipeline(stages=[labelIndexer, featureIndexer, gbt])

    # Train model. This also runs the indexers.
    model = pipeline.fit(trainingData)

    # Save the model
    model.write().overwrite().save(model_path)

    ##### Make predictions with test Data
    print('TEST DATA')

    # Make predictions.
    predictions = model.transform(testData)

    # Select example rows to display.
    predictions.select("prediction", "indexedLabel", "features").show(5)
```



```

# Select (prediction, true label) and compute test error
evaluator = MulticlassClassificationEvaluator(
    labelCol="indexedLabel", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Test Error = %g" % (1.0 - accuracy))

gbtModel = model.stages[2]
print(gbtModel) # summary only

##### Make predictions with training Data
print('TRAINING DATA')

# Make predictions.
predictions = model.transform(trainingData)

# Select example rows to display.
predictions.select("prediction", "indexedLabel", "features").show(5)

# Select (prediction, true label) and compute test error
evaluator = MulticlassClassificationEvaluator(
    labelCol="indexedLabel", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Test Error = %g" % (1.0 - accuracy))

gbtModel = model.stages[2]
print(gbtModel) # summary only

return testData, trainingData

read_path = '../Data/'
file_name = 'data_model_libsvm.csv'
libsvm_filename = read_path + file_name
data = spark.read.format("libsvm").load(libsvm_filename)
test, train = gradient_boosted_tree(data)
112         else: # must be an estimator
--> 113             model = stage.fit(dataset)
114             transformers.append(model)
115             if i < indexOfLastEstimator:

/home/nacho/spark/python/pyspark/ml/base.py in fit(self, dataset, params)
62         return self.copy(params)._fit(dataset)
63     else:

```

```
---> 64             return self._fit(dataset)
      65         else:
      66             raise ValueError("Params must be either a param map or a list/tuple of param maps, ")

/home/nacho/spark/python/pyspark/ml/wrapper.py in _fit(self, dataset)
      211
      212     def _fit(self, dataset):
--> 213         java_model = self._fit_java(dataset)
      214         return self._create_model(java_model)
      215

/home/nacho/spark/python/pyspark/ml/wrapper.py in _fit_java(self, dataset)
```

Multilayer perceptron classifier

```
In [3]: from pyspark.ml.classification import MultilayerPerceptronClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

MPCmodel_path = "../Model/Multi_percep"

def multilayer_perceptron(data, layers, model_path = MPCmodel_path):

    # Split the data into train and test
    splits = data.randomSplit([0.6, 0.4], 1234)
    train = splits[0]
    test = splits[1]

    # create the trainer and set its parameters
    trainer = MultilayerPerceptronClassifier(maxIter=100, layers=layers, blockSize=128, seed=1234)

    # train the model
    model = trainer.fit(train)

    # Save the model
    model.write().overwrite().save(model_path)

    ##### Make predictions with test Data
    print('TEST DATA')
    # compute accuracy on the test set
    result = model.transform(test)
    predictionAndLabels = result.select("prediction", "label")
    evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
    print("Test set accuracy = " + str(evaluator.evaluate(predictionAndLabels)))

    ##### Make predictions with test Data
    print('TRAINING DATA')
    # compute accuracy on the test set
    result = model.transform(train)
    predictionAndLabels = result.select("prediction", "label")
    evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
    print("Training set accuracy = " + str(evaluator.evaluate(predictionAndLabels)))

    return test, train
```

```
In [4]: read_path = '../Data/'
        libsvm_filename = read_path + 'data_model_libsvm_multi.csv'

        data = spark.read.format("libsvm").load(libsvm_filename)

        # specify layers for the neural network:
        # input layer of size 2 (features), two intermediate of size 5 and 4
        # and output of size 3 (classes)
        layers = [2, 5, 4, 3]

        test, train = multilayer_perceptron(data, layers)

TEST DATA
Test set accuracy = 0.5251414795944537
TRAINING DATA
Training set accuracy = 0.5261753046875888
```

```
In [5]: read_path = '../Data/'
        libsvm_filename = read_path + 'data_model_libsvm_multi.csv'

        data = spark.read.format("libsvm").load(libsvm_filename)

        # specify layers for the neural network:
        # input layer of size 2 (features), several intermediate
        # and output of size 3 (classes)
        layers = [2, 5, 4, 4, 7, 3]

        test, train = multilayer_perceptron(data, layers)

TEST DATA
Test set accuracy = 0.5177212809246183
TRAINING DATA
Training set accuracy = 0.5191331311548029
```

Random Forest

General example with all data

```

In [30]: RFmodel_path = "../Model/RandomForest_Batch"
def randomforest(data, model_path = RFmodel_path):
    # Index labels, adding metadata to the label column.
    # Fit on whole dataset to include all labels in index.
    labelIndexer = StringIndexer(inputCol="label", outputCol="indexedLabel").fit(data)

    # Index labels, adding metadata to the label column.
    # Fit on whole dataset to include all labels in index.
    # labelIndexer = StringIndexer(inputCol="label", outputCol="indexedLabel").fit(data)

    # Automatically identify categorical features, and index them.
    # Set maxCategories so features with > 4 distinct values are treated as continuous.
    featureIndexer = \
        VectorIndexer(inputCol="features", outputCol="indexedFeatures", maxCategories=4).fit(data)

    # Split the data into training and test sets (30% held out for testing)
    (trainingData, testData) = data.randomSplit([0.7, 0.3])

    # Train a RandomForest model.
    rf = RandomForestClassifier(labelCol="indexedLabel", featuresCol="indexedFeatures", numTrees=100)

    # Convert indexed labels back to original labels.
    labelConverter = IndexToString(inputCol="prediction", outputCol="predictedLabel",
                                   labels=labelIndexer.labels)

    # Chain indexers and forest in a Pipeline
    pipeline = Pipeline(stages=[labelIndexer, featureIndexer, rf, labelConverter])

    # Train model. This also runs the indexers.
    model = pipeline.fit(trainingData)

    # Save the model
    model.write().overwrite().save(model_path)

    ##### Make predictions with test Data
    print('TEST DATA')
    predictions = model.transform(testData)

    # Select example rows to display.
    predictions.select("predictedLabel", "label", "features").show(5)

```

```

# Select (prediction, true label) and compute test error
evaluator = MulticlassClassificationEvaluator(
    labelCol="indexedLabel", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Test Error = %g" % (1.0 - accuracy))

rfModel = model.stages[2]
print(rfModel) # summary only
print('\n\n\n')

##### Make predictions with train Data
print('TRAINING DATA')
predictions = model.transform(trainingData)

# Select example rows to display.
predictions.select("predictedLabel", "label", "features").show(5)

# Select (prediction, true label) and compute test error
evaluator = MulticlassClassificationEvaluator(
    labelCol="indexedLabel", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Test Error = %g" % (1.0 - accuracy))

rfModel = model.stages[2]
print(rfModel) # summary only
return testData, trainingData

read_path = '../Data/'
file_name = 'data_model_libsvm.csv'
libsvm_filename = read_path + file_name
data = spark.read.format("libsvm").load(libsvm_filename)
test, train = randomforest(data)

```

TEST DATA

```

+-----+-----+-----+
|predictedLabel|label|      features|
+-----+-----+-----+
|           1.0|  1.0|(2,[0],[1.0])|
|           1.0|  1.0|(2,[0],[1.0])|
|           1.0|  1.0|(2,[0],[1.0])|
|           1.0|  1.0|(2,[0],[1.0])|
|           1.0|  1.0|(2,[0],[1.0])|

```

```
+-----+-----+-----+
```

only showing top 5 rows

Test Error = 0.215483

RandomForestClassificationModel (uid=rfc_053eced0c34b) with 100 trees

TRAINING DATA

```
+-----+-----+-----+
```

```
|predictedLabel|label|      features|
```

```
+-----+-----+-----+
```

```
|          1.0|  1.0|(2,[0],[1.0])|
```

```
|          1.0|  1.0|(2,[0],[1.0])|
```

```
|          1.0|  1.0|(2,[0],[1.0])|
```

```
|          1.0|  1.0|(2,[0],[1.0])|
```

```
|          1.0|  1.0|(2,[0],[1.0])|
```

```
+-----+-----+-----+
```

only showing top 5 rows

Test Error = 0.211873

RandomForestClassificationModel (uid=rfc_053eced0c34b) with 100 trees

Predicting only one value

```

In [9]: from pyspark.mllib.tree import RandomForest, RandomForestModel
        from pyspark.mllib.util import MLUtils

        RF_streaming_path = '../Model/RandomForest_Streaming'

        NUM_TREES = 100

        read_path = '../Data/'
        file_name = 'data_model_libsvm.csv'
        libsvm_filename = read_path + file_name

        # Load and parse the data file into an RDD of LabeledPoint.
        data = MLUtils.loadLibSVMFile(sc, libsvm_filename)
        # Split the data into training and test sets (30% held out for testing)
        (trainingData, testData) = data.randomSplit([0.7, 0.3])

        # Train a RandomForest model.
        # Empty categoricalFeaturesInfo indicates all features are continuous.
        # Note: Use larger numTrees in practice.
        # Setting featureSubsetStrategy="auto" lets the algorithm choose.
        model = RandomForest.trainClassifier(trainingData, numClasses=4, categoricalFeaturesInfo={},
                                             numTrees=NUM_TREES, featureSubsetStrategy="auto",
                                             impurity='gini', maxDepth=4, maxBins=32)

        # Evaluate model on test instances and compute test error
        predictions = model.predict(testData.map(lambda x: x.features))

        labelsAndPredictions = testData.map(lambda lp: lp.label).zip(predictions)
        testErr = labelsAndPredictions.filter(
            lambda lp: lp[0] != lp[1]).count() / float(testData.count())

        print('Test Error = ' + str(testErr))
        print('Learned classification forest model:')
        print(model.toDebugString())

        # Save and load model
        RF_streaming_path = '../Model/RandomForest_Streaming'
        if os.path.isdir(RF_streaming_path):
            # If exists
            shutil.rmtree(RF_streaming_path)

```



```
# Save
model.save(sc, RF_streaming_path)
# Load
sameModel = RandomForestModel.load(sc, RF_streaming_path)
```

Test Error = 0.21548003089645146
Learned classification forest model:
TreeEnsembleModel classifier with 100 trees

Tree 0:
If (feature 0 <= 2.0)
If (feature 1 <= 4344.0)
If (feature 1 <= 1418.0)
If (feature 1 <= 0.0)
Predict: 1.0
Else (feature 1 > 0.0)
Predict: 3.0
Else (feature 1 > 1418.0)
If (feature 1 <= 2896.0)
Predict: 1.0
Else (feature 1 > 2896.0)
Predict: 1.0
Else (feature 1 > 4344.0)
If (feature 1 <= 7240.0)
If (feature 1 <= 5701.0)

In []: label_keys

In []: proto_keys

In []: value = [1, 1090]
sameModel.predict(value)

In []: *# For predicting just running this chunk*
from pyspark.mllib.tree **import** RandomForest, RandomForestModel
from pyspark.mllib.util **import** MLUtils
sameModel = RandomForestModel.load(sc, RF_streaming_path)
value = [3, 175]
sameModel.predict(value)

In [8]:

In []: